

# Adapting CheckList to ELECTRA for QA: Fine-Tuning for Challenge Datasets

Anonymous CS388 submission

## Abstract

Several challenges have been made to the SQuAD dataset over time. These include issues with robustness, biases in the training data, and lack of “difficult” or multi-step question answering. Challenge datasets such as Adversarial SQuAD (Jia and Liang, 2017) and Adversarial QA (Bartolo et. al, 2020) have been proposed to strengthen Question Answering models. Additionally, testing methods such as CheckList (Ribeiro et. al, 2020) have been designed to expose weaknesses in existing datasets and testing methodologies for QA models. The goals of this paper are three-fold: first, to adapt CheckList tests to a SQuAD-trained ELECTRA QA model (Clark et. al, 2019) to understand its weaknesses and strengths. Second, to train it using adversarial datasets and evaluate which categories improved. Lastly, attempting to use hand-tuned training sets to further strengthen performance on specific types of tests from CheckList. I will show that while it is possible to improve performance using hand-tuned sets on certain categories, other categories may also see either increased or reduced performance. Even though a “one-size-fits-all” QA model was not achieved in this experiment, high performance on specific types of questions is demonstrated.

## 1 Introduction

Question Answering (QA), while historically a task given to computers as early as the 1960s with a Baseball-specific question/answer machine (Green et. al, 1961),<sup>1</sup> it has come into the mainstream within the last decade, and the SQuAD dataset was introduced at Stanford in 2016 (Rajpurkar et. al, 2016) to enable widespread development and training of QA models on over 100,000 labeled examples.

<sup>1</sup> In hindsight it seems obvious that baseball statistics nerds would be one of the earliest to design a computer system to query their mountains of stats.

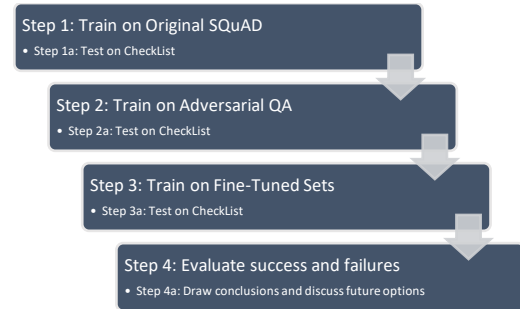


Figure 1: Illustration of general flow of experiment, utilizing adversarial datasets like those of Jia and Liang (2017) and inspired by fine-tuning methods of Liu et. al (2019).

However, over time, challenges have been introduced to models developed and trained using open-source datasets like SQuAD, exposing weaknesses and areas with room for improvement (Jia and Lang, 2017; Rajpurkar et. al, 2018; Bartolo et. al, 2020). An important part of understanding and improving models such as those used for Question Answering is knowing the limitations of the models: what they can do, and more importantly what they cannot do. For example, SQuAD-trained QA models have not been shown to demonstrate “multi-hop” question answering (Chen and Durrett, 2019), as questions in the dataset do not demand “multi-hop” reasoning such as those in the HotpotQA dataset (Yang et. al, 2018).

I studied the literature referenced above and others, and chose to train the ELECTRA model using the Huggingface Transformers framework (Wolf et. al, 2019), utilizing starter code provided in the assignment template. Initial training was done using the ‘train’ split of the SQuAD dataset. Evaluation was performed by adapting the

62 CheckList tests to work with the provided code  
 63 and HuggingFace Transformer framework.  
 64 CheckList suggests various test categories,  
 65 including Vocabulary, Taxonomy, Robustness,  
 66 Temporal, Negation, Coreference, and Semantic  
 67 Role Labeling (SRL); each is evaluated using  
 68 percentage of tests failed (presented in the paper  
 69 later on as percentage of tests passed). Further  
 70 training was conducted by introducing data from  
 71 adversarial QA challenge sets to inoculate the  
 72 model (Liu et. al 2019), such as AddOneSent and  
 73 AddSent from the Adversarial Examples paper’s  
 74 squad\_adversarial dataset (Jia and Lang,  
 75 2017) and adversarial\_qa dataset (Bartolo et.  
 76 al, 2020). I initially hypothesized improvements  
 77 could be made to CheckList test category  
 78 performance using only samples of adversarial  
 79 QA data for training, but this was not consistently  
 80 backed by results. This motivated me to pursue  
 81 fine-tuning training using small datasets to  
 82 improve performance on specific test types.  
 83 I compared the percentage performance of  
 84 each testing category from CheckList over  
 85 various combinations of SQuAD and adversarial  
 86 data attempts and attempted to train the model on  
 87 specific hand-made datasets to help the model  
 88 learn the concepts of negation and former/latter  
 89 relationships. Both categories were routinely  
 90 scoring approximately 0.0% regardless of  
 91 underlying dataset and training. Training data  
 92 was adapted from the CheckList tests, but  
 93 specifically randomizing sentence structure,  
 94 names, professions, and answers and keeping the  
 95 negation and former/latter relationships. This  
 96 data, after one training pass, increased  
 97 performance from 0% passing to upwards of  
 98 60%, as well as incidental improvements on other  
 99 categories of CheckList tests. In most cases of  
 100 adversarial and small dataset training, F1  
 101 performance on the standard SQuAD dataset  
 102 was improved from my initial baseline, and in all  
 103 cases never significantly diminished in  
 104 performance. Finally, I also evaluated the  
 105 various trained models on the adversarial\_qa  
 106 model from Bartolo et. al, 2020 and saw  
 107 significant improvement from the base SQuAD-  
 108 trained ELECTRA model, but only an increase to  
 109 approximately 40% F1; nothing close to state-of-  
 110 the-art performance.

---

<sup>2</sup>This option will not be available after “mid-December”  
 2022, according to Google Colab documentation:  
<https://github.com/googlecolab/colabtools/issues/3246>

111 I show in this paper that significant  
 112 performance improvements can be made to  
 113 specific query types using fine-tuning and  
 114 CheckList on an ELECTRA model for Question  
 115 Answering, with minimal loss of generalization  
 116 for other types of queries and minimal overall  
 117 performance impact on general SQuAD F1  
 118 performance.  
 119

## 120 2 Methodology

### 121 2.1 Initial Setup/Hardware Configuration

122 This experiment was conducted primarily on  
 123 Google Colab with “GPU” and “premium GPU”  
 124 options, in order to maximize performance  
 125 during computation-intensive training and  
 126 evaluation phases. Colab Pro+ subscription was  
 127 utilized to minimize disconnect from runtime,  
 128 which would have negatively impacted the  
 129 performance of the experiments. Python version  
 130 3.7.15 was used due to an incompatibility  
 131 between the default Colab Python 3.8  
 132 environment and the CheckList testing suite; this  
 133 downgrade was achieved by selecting “Use  
 134 fallback runtime version” from the Colab  
 135 Command Palette.<sup>2</sup>

### 136 2.2 Initial Training and Evaluation

137 I started with a “stock” ELECTRA model from  
 138 the HuggingFace Transformers pipeline and  
 139 trained it using default parameters on the SQuAD  
 140 database to achieve my “baseline” results. The  
 141 first evaluation was conducted by evaluating the  
 142 model on exact match and F1; receiving 78.01%  
 143 and 86.09, respectively.

144 From there, I downloaded and modified the  
 145 CheckList testing package from  
 146 <https://github.com/marcotcr/checklis>  
 147 t. I modified the SQuAD loading starter code  
 148 provided in the CheckList repository in order to  
 149 properly format the SQuAD test package as a  
 150 JSONL datafile that Huggingface Datasets could  
 151 easily import with existing code (see Appendix A  
 152 for function details). My initial plan was to  
 153 develop my own tests using the CheckList  
 154 categories but after spending multiple days  
 155 attempting to understand and implement these  
 156 tests using the included packages, unfamiliarity

with Python and Huggingface Pipelines prevented me from making full use of the CheckList suite (including Editor and TestSuite functions). After significant effort was expended, I changed course and ended up falling back on the provided TestSuite which runs over 79,000 prebuilt tests on the categories in the CheckList paper (Ribeiro et. al, 2020).

The baseline SQuAD model performed well on certain CheckList suite test categories and failed significantly on many others; the details will be shown in the “Results” section below. The subsequent training was based off either the “stock” ELECTRA model or this SQuAD-trained model, referred to “baseline model,” or “baseline” in the following sections.

### 2.3 Subsequent Testing and Evaluation Methodology – squad\_adversarial

My next goal was to run several experiments to determine how the ELECTRA model could be enhanced to perform better on CheckList suite tests. I ran experiments to train the model on two different Adversarial QA datasets: squad\_adversarial (Jia and Lang, 2017) and adversarial\_qa (Bartolo et. al, 2020) under various subsets and combinations of these datasets.

First, I trained the “stock” model on the AddOneSent subset of the squad\_adversarial dataset combined with SQuAD dataset via Huggingface `Datasets.combine_dataset()` functionality and noted slight improvement on exact match % and F1, as well as improvement on Taxonomy type tests from CheckList but minimal improvement or slight decreases on others.

Next, I trained the stock model on the AddSent subset of the squad\_adversarial dataset combining with SQuAD dataset; the goal was to introduce the model to more examples of adversarial QA data in hopes of improving performance on SQuAD data and CheckList tests, as AddSent contains more examples of adversarial sentences. This resulted in another slight improvement to exact match % and F1 on SQuAD and adversarial\_qa validation sets, as well as subsequent improvements on some Robustness and Taxonomy tests, but no significant improvements or changes on most

other types of tests. I decided I needed to investigate a second adversarial QA dataset.

### 2.4 Testing and Evaluation Methodology – adversarial\_qa

From there, I investigated the adversarial\_qa dataset created by Bartolo et. al 2020, and decided it might add some more robustness to the model. I conducted two different experiments with this dataset; first performing subsequent training on the “baseline” SQuAD-trained model with this challenge set, as well as using combine\_dataset to combine this dataset with AddSent of squad\_adversarial dataset.

Conducting follow-on training of the baseline model with challenge data, as opposed to combining the challenge data with the SQuAD dataset to train the “stock” model, showed some surprising results. This includes improving significantly on certain CheckList sections like comparison and taxonomy, and the ability to distinguish “his/her” relationships at a 48% rate, but also a significant performance drop in categories that other experiments retained or improved, including Robustness, Named Entity Recognition, and Fairness. Due to the uneven results, this model of “follow-on” training of the stock model was not continued for my subsequent experiments.

The next experiment was combining both adversarial challenge datasets with the original SQuAD dataset and conducting full training cycle on this combined “triple” dataset. This dataset showed the strongest performance of all experiments on exact match % and F1; 80.49% and 87.54, respectively. It also showed similar or improved performance on nearly every CheckList test category to the baseline SQuAD-trained model, with only one category decreasing more than five percentage points (his/her distinction). Due to these results, I decided to adopt this “triple” trained model as the new baseline; subsequent references to this model will be noted as “triple” or “triple dataset” model.

### 2.5 Testing and Evaluation – Fine-Tuning

Taking the “triple” dataset trained model, I was motivated to investigate whether or not specific categories of CheckList tests could be approached by training the model further on a small fine-tuned dataset with specific examples.

Specifically, my goal was to train the model on the relationship between “former/latter” and named entities earlier in the passage, and the concept of negation of earlier statements/descriptors. I approached each one separately, measured performance on overall SQuAD data and CheckList tests, and then attempted training with a combined fine-tuned dataset with both categories.

I noted that previous models showed close to 0.0% accuracy on the Former / Latter relationship portion of CheckList test suite and, motivated by the Inoculation by Fine-Tuning paper (Liu et. al, 2019), I decided to adapt small training sets (approximately 100 entries) with explicit negation in both the context and the question to evaluate whether the model could successfully “learn” this relationship. I also created a dataset of approximately 100 entries with explicit former / latter relationships. In both fine-tuned datasets, I randomized names, adjectives, descriptors, etc. using the following sites: in an attempt to prevent the model from learning specific words as features when the tests were run. I first performed follow-on training with this small former/latter dataset on the “baseline” SQuAD trained model, noting an increase from 0.0% of tests passing to only 6.1% passing, but no significant decreases in performance on overall SQuAD data or other CheckList categories. Due to these results, I then decided to train the “triple” trained model with the former/latter dataset and noted marked improvements in both former/latter recognition (42.95% accuracy) and significant progress in Taxonomy and Fairness category tests, as well.

I then attempted to train the model on the former-latter dataset for a second time, but this resulted in slight decreases across the board, including even a slight decrease in “former/latter” test performance so I did not pursue this approach. The success on the single training cycle of “former/latter” fine-tune dataset motivated me to create the negation dataset mentioned above and evaluate the results of training the “triple” model on another small dataset to fine-tune results.

I created the dataset of approximately 125 examples to train the model on examples of negation: specifically including passages with negation and questions with negation, as well as samples with negation only mentioned in the question. I found that the triple dataset evaluated

very strongly on this category after just training on a small set of examples, as high as 99.38% accuracy on this subset of tests. While I did randomize names, descriptors, professions, etc., to reduce chance of over-fitting the model it is possible that the model was still able to easily over-fit this training data and I will discuss this possibility in results and conclusions below. After evaluating the model on negation, I wanted to combine the two datasets and attempt to train the model on both concepts: negation and former/latter.

I then combined the two datasets and trained the “triple” model on a former/latter and negation fine-tuning set and evaluated results on SQuAD and CheckList. Overall, the fine-tuned model performed similarly to the “triple” and “baseline” models on overall SQuAD performance, losing only about 1 point off F1 from the triple dataset and still 1 point higher than the baseline, but showed significant improvement in several categories of CheckList tests, not just negation and former/latter. Specific results and conclusions as to why will be drawn later in this paper.

## 2.6 Resources – Scripts, Notebook, Datasets, Results

All datasets used, training scripts modified, and the Jupyter notebook I ran in Google Colab, as well as results received are listed below:

<https://anonymfile.com/y6kBy/results.xlsx>  
<https://anonymfile.com/OaX2Q/run.py>  
<https://anonymfile.com/lyPBV/requirements.txt>  
<https://anonymfile.com/KV2Q7/nlp-final.ipynb>  
<https://anonymfile.com/0mPob/f1-eval-results-for-various-models.txt>  
<https://anonymfile.com/znmaj/helpers.py>  
<https://anonymfile.com/DXNm9/squad-negation-2.jsonl>  
<https://anonymfile.com/xJroy/squad-negationformerlatter.jsonl>  
<https://anonymfile.com/2pWa1/squad-formerlattertrain.jsonl>  
<https://anonymfile.com/n051a/squad-formatted.jsonl>  
<https://randomwordgenerator.com/name.php>  
<https://www.randomready.com/random-job-generator/>



**Results by Category on SQuAD Validation Sets and Selected CheckList Tests**

Training Datasets	SQuAD validation (exact % / F1)	Adversarial QA validation (exact % / F1)	Vocab Comparison/ Synonyms	Taxonomy attributes/ nationality- job/animal- vehicle	Fairness Male- Female Stereo- types	Temporal Change in job	Negation	Coreference His-her	Former -Latter
SQuAD	78.01%/86.09	19.93%/30.09	1.21%/97.09%	1.40%/35.40% / 36.35%	56.52%	99.59%	6.02%	6.20%	0.00%
SQuAD + AddSent	80.24%/87.32	18.07%/28.01	0.81%/98.88%	1.00%/63.20%/ 38.86%	8.00%	84.23%	0.00%	0.20%	0.00%
SQuAD then AddSent	69.92%/78.54	<b>29.23%/39.61</b>	45.95%/84.79%	<b>69.20%/79.40%</b> /52.71%	4.60%	98.76%	0.81%	<b>48.20%</b>	0.00%
SQuAD + adversarial + AddSent (Triple)	<b>80.49%/87.54</b>	28.47%/39.00	2.02%/95.75%	10.80%/40.80% /48.09%	55.10%	96.47%	4.49%	1.00%	0.00%
SQuAD then former-latter	77.23%/84.69	21.03%/29.26	1.01%/92.39%	4.40%/32.60%/ 34.34%	6.95%	8.30%	0.20%	37.60%	6.11%
Triple then former-latter set	79.60%/86.39	28.00%/37.87	10.12%/44.07%	37.40%/47.00% /37.25%	25.25%	6.85%	0.00%	39.80%	42.95%
Triple then negation set	80.04%/87.00	28.30%/38.21	64.78%/97.32%	6.20%/19.80%/ <b>74.20%</b>	<b>100.00%</b>	<b>100.00%</b>	<b>99.38%</b>	0.06%	0.00%
Triple then negation+ former-latter set	79.46%/86.58	28.00%/37.34	<b>78.34%/97.76%</b>	16.60%/6.20%/ 70.08%	<b>100.00%</b>	99.59%	<b>99.79%</b>	4.40%	<b>60.21%</b>

Table 1: Performance of experimental models on SQuAD and various CheckList test categories. Not every CheckList category is included due to spacing; see the full spreadsheet at <https://anonymfile.com/y6kBy/results.xlsx> for more details.

Note that most benchmarks are achieved on either the SQuAD-trained “baseline” followed by subsequent training on adversarial data, or on the “triple” set trained further on negation/former-latter datasets.

### 3 Results

Selected results are shown in Table 1; all else in the spreadsheet hyperlinked in Table 1’s caption.

#### 3.1 Results from Baseline SQuAD training

Training the stock ELECTRA model on SQuAD ‘train’ split dataset resulted in results of 78.01% exact match and 86.09 F1. Evaluation on the ‘validation’ set of the adversarial\_qa dataset (Bartolo et. al, 2020) resulted in only 19.93% exact match and 30.09 F1. This closely reproduces the results of Jia and Lang (2017) in that strong performance on the standard SQuAD is not nearly robust enough to perform well on human-created adversarial examples.

Results from Baseline model on CheckList tests enabled several observations, as well. The initial SQuAD dataset enables the model to pass the Vocabulary test category of “Synonyms” well, as this is well-represented in the question/answer set. I achieved 97.09% passing on CheckList Synonyms just with the baseline SQuAD training data. However, the next two Vocabulary tests, Comparison and Intensifiers failed, passing only 1.21% and 0.00% of tests, respectively. This motivated me to examine the

SQuAD dataset and I found that neither of these two examples were represented well in the underlying data. While comparison is something that performed better upon using fine-tuning datasets later on, the datasets I used did not improve performance on Intensifiers and I would have liked to investigate separately if I had more time.

Baseline SQuAD model showed mixed results on Taxonomy category; failing Size/Shape/Age/Color tests with 1.40% passing, and reaching only 35.40% and 36.35% respectively for distinguishing Nationality from Profession and distinguishing Animals from Vehicles. These categories showed drastic improvement when training on challenge data, discussed later. The baseline model also failed on Antonym tests and Coreference tests, but none of the models achieved significant success on these categories.

Additionally, the SQuAD-trained baseline model performed very well on Robustness tests and Named Entity Recognition, passing Question Typo with 75.20%, Question Contractions with 88.00%, Adding Random Sentence with 84.60% (this one specifically was the highest of all models). Changing Name and Changing

Location tests passed with 91.60% and 84.20% respectively.

Mixed results were achieved on Temporal category tests, with the baseline model passing 99.59% of “Change in Profession” temporal tests but passing 0.00% of “before/after to first/last” tests. None of the models achieved success on the “before/after to first/last” categories and this is also something I would investigate if fine-tuning could help if I had more time.

Semantic Role Labeling was something that the baseline model, as well as subsequent models, largely failed to perform on and deserves further investigation.

Finally, the baseline model achieved poor performance on Negation tests and Former / Latter tests. The baseline model achieved 6.02% on Negation tests and 0.00% on Former / Latter tests; this was part of my motivation to fine-tune on these test cases.

### 3.2 Results from Adversarial training

Training on various adversarial datasets, including those from [Jia and Lang \(2017\)](#) and [Bartolo et. al \(2020\)](#) resulted in insignificant changes in performance when only `squad_adversarial` data was combined with original SQuAD data, but showed significant results when the model was trained with `squad_adversarial` as a second set following initial baseline SQuAD training, as well as when both `squad_adversarial` and `adversarial_qa` datasets were combined together with SQuAD dataset.

Overall, when `AddSent` or `AddOneSent` were inserted into the SQuAD dataset, the F1 performance on both SQuAD and `adversarial_qa` validation sets were minimally impacted; less than 1.5 points each. `CheckList` tests were minimally impacted as well; however, ability to distinguish Nationality from Profession in the Taxonomy category did increase from 35.40% in the baseline to 63.20% with `AddSent` data included. It became clear that just training with adversarial data mixed into SQuAD data would not be enough to make significant progress on either the SQuAD F1, the `adversarial_qa` F1, or `CheckList` test categories. This conclusion motivated training with `squad_adversarial` dataset as a secondary training set instead of mixed-in to see if values could be optimized further from SQuAD baseline training.

I then trained the baseline SQuAD-trained model on the `squad_adversarial` `AddSent` dataset as a secondary dataset. This resulted in a sharp drop in F1 performance on SQuAD data, similar to that shown by [Liu et. al in 2019](#), but in fact significant performance changes in evaluation on `adversarial_qa` validation set (39.61 F1, the highest achieved of any model tested), and several `CheckList` test categories. First, the model gained approximately 45% ability to distinguish Comparison tests, and succeeded the best of any model tested on several Taxonomy tests, including Size/Shape/Age/Color and Profession vs. Nationality; 69.20% and 79.40% each, respectively. Coreference performance improved drastically, increasing from 6.20% to 48.20% as well. This model trained on adversarial data following SQuAD data did not fare as well on Male/Female stereotypes, Robustness tests, or Named Entity Recognition tests; but overall it did not lose more than approximately 10-15 percentage points on any; this approach shows promise for those categories (Taxonomy and Comparison, as well as Coreference) if one is willing to take the performance hit on other categories. The overall performance on SQuAD decreasing by about 9 points F1 meant that I did not utilize this method in further experimentation, but I would certainly investigate more if I had more time or larger scope.

After abandoning separate dataset sequential training, I then combined both adversarial sets with SQuAD data and trained the stock model on this combined dataset, referred to earlier and below as the “triple” dataset. This achieved the highest performance on exact match and F1 (80.49% and 87.54 respectively), as well as second highest F1 on `adversarial_qa` set (39.00 vs the previous method’s 39.61). Additionally, this model did not lose more than 1-3 percentage points of performance on any `CheckList` category from the baseline SQuAD trained model, and gained anywhere between 5-15% performance increases on many `CheckList` categories. This increased performance may be attributable to the variance in questions and variance in passages utilized by combining different data sources. I chose to continue with this as the primary model moving forward due to the promising performance shown. From here,

my goal was to approach specific CheckList test categories and attempt to improve performance on those questions without loss of generality on SQuAD, adversarial\_qa, and the other CheckList test categories.

### 3.3 Results from Fine-Tuning training

Overall results show that performance can be sharply increased on specific test categories using small (~100 entry) datasets fine-tuned to the task given, without large general loss in overall performance or on any given category of specific CheckList test.

As described in methodology, the first attempt was training the baseline SQuAD model on a former-latter challenge set; this showed slight decreases in generalized performance (exact match and F1), as well as on various CheckList test categories, such as Fairness and Temporal decreasing from 56.52% -> 6.95% and 99.59% -> 8.30%, respectively. It only significantly gained in Semantic Role Labeling (Agent/Object), increasing to 19.32% from 9.46% as well as Coreference His/Her, increasing to 37.60% from 6.20%. Even Former/Latter relationship tests only increased to 6.11% from the baseline result of 0.00%. This was discouraging to see, but I was motivated to try this dataset on the “triple” trained model.

When trained further with fine-tuning sets like former-latter and negation, the “triple” trained model performed strongly on both general performance and specific categories. For example, the triple-trained model with former-latter training achieved 42.95% passing on the former/latter test, but also significant increases in performance in Coreference his/her, Taxonomy tests. Performance did decrease on Fairness and Temporal; this merited further investigation and in fact performance was restored when training with the addition of the “negation” dataset.

Training the triple model on the negation challenge set showed promising results, sharply increasing performance on Comparison tests, Animal vs. Vehicle Taxonomy, Male/Female stereotypes, Change in profession, and Negation tests (passing the last item with 99.38%). Such a high rate of passing might be attributed to overfitting; I tried to minimize this by randomizing professions, descriptors, names, and other words in the tests but due to the overall structure remaining similar it may not generalize

to larger passages. Since both negation and former/latter relationships require locality, though (the two sentences with the relationship will still be next to each other, in the same relative order), I am optimistic that it would still apply in longer, more general passages. My next step in a further experiment would be to analyze this performance further to determine if the data was overfit, and if so, how general performance on these types of questions (beyond CheckList) fares.

Overall, combining the two fine-tuning sets (Negation set and Former/Latter set) and performing additional training with the combination on the “triple” trained model showed benchmark or near-benchmark performance on several categories: Vocabulary Comparison, Comparison-Antonyms, Male/Female Stereotypes, Negation, Temporal Change in Profession, Former / Latter (60.21%), and was within 1-3 points of the majority of every other CheckList category. Additionally, it was still ahead in F1 performance on both the SQuAD validation set and the adversarial\_qa validation set. It remains to be seen whether all these specific results can be generalized to larger datasets, but I can confidently state that test categories from CheckList can be trained with fine-tuning without loss of generality on overall performance.

### 3.4 Discussion

As mentioned above multiple times, it is not entirely clear that training on the fine-tuned datasets will generalize to larger datasets. Investigating whether other categories of CheckList tests can be fine-tuned as well as attempting to test the model on more varied passages to evaluate CheckList category performance would be the next logical step for this experiment, but it would be out of scope (and time) for a single person project in this context.

The fact that performance benchmarks were reached in some categories using both adversarial training data and fine-tuning data speaks to the promise of each method. I also draw the conclusion that a QA model can be adapted to specific types of questions more successfully; if a QA model is going to be used to answer Taxonomy type questions, it may be worth a tradeoff of loss of generality on other categories to improve performance on that specific type of

question. Perhaps a higher model could even be trained to determine type of question asked, which could then be fed to a lower QA model specifically trained on those types of questions/features. This idea is far out of the scope of this project, but is an interesting thought experiment.

Another conclusion drawn via my observations is that the “Male/Female stereotype” test was hugely influenced by training data and approach. This helped me realize how sensitive these types of models are to biased input, and the importance of those working on commercial models being aware of the propensity for bias in complex neural models. If models are being updated consistently with new data/search queries, the teams must be equally watchful for bias or decreased performance on certain categories occurring.

Finally, none of the training I did impacted performance on Intensifiers, asking questions with opposite Antonyms, “He/She” coreferences, or Semantic Role Labeling with three agents/objects in the sentence. I would pursue fine-tuning on these test categories, but the fact that there was not any incidental improvement from adversarial data or fine-tuning other categories, as many other CheckList test categories showed, might mean that these types of linguistic patterns are more difficult to adapt to an ELECTRA SQuAD-trained model.

I remain optimistic that Former/Latter and Negation features will be similar in larger more varied datasets than those evaluated in the CheckList release package, as they are very specific structures and orders of sentences, but I would be very interested to develop more complex context/question pairs to test these categories in a further project and rule out over-fitting.

Overall, these results demonstrate that CheckList can be a valuable tool to analyzing performance of a Question-Answering model, and that specific types of bugs can be fixed or performance improved by using Fine-Tuning to teach the model those features.

## 4 Conclusion

Training an ELECTRA QA model on SQuAD followed by adversarial data and fine-tuning using small datasets can illuminate performance issues with current QA models and datasets.

While “passing” performance was not achieved on an adversarial dataset, significant improvement was shown in F1 score (from 30.09 to as high as 39.61). Improvement was made on several categories of test from the CheckList suite, without loss of performance on the vast majority of CheckList category or significant decrease to overall F1 performance on SQuAD validation set. A few categories of CheckList tests were not affected by any adversarial training or fine-tuning; it remains to be seen if those categories can be fine-tuned or if they are a more complex challenge than a SQuAD-trained ELECTRA model can address. The incidental performance increase in other categories of CheckList tests may be attributed to features of the CheckList suite questions being too similar between categories, or due to over-fitting of data. Nevertheless, the methods show promise and deserve evaluation over more complex testing datasets, likely derived from CheckList or another evaluation method.

Additionally, male/female and other built-in “biases” are an important topic that should continue to be researched further to avoid potential negative externalities of these types of biases as “AI” and “Machine Learning” continue to become buzzwords and complex models that may not be fully understood by those using them get implemented into everyday systems like job hiring processes.

Overall, the process of training on adversarial data and fine-tuning on small datasets shows promise for improving performance of SQuAD-trained ELECTRA models and potentially other types of complex Question Answering models.

## Acknowledgments

I would like to acknowledge my wife, who actually has a published scientific paper to her name ([Theisen et. al, 2017](#)), and thus indulged me by making me feel more important about this project/report than I maybe should have.



## References

- Max Bartolo, Alastair Roberts, Johannes Welbl, Sebastian Riedel, and Pontus Stenetorp. 2020. [Beat the AI: Investigating Adversarial Human Annotation for Reading Comprehension](#). *Transactions of the Association for Computational Linguistics*, 8:662–678.
- Jifan Chen and Greg Durrett. 2019. [Understanding Dataset Design Choices for Multi-hop Reasoning](#). In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 4026–4032, Minneapolis, Minnesota. Association for Computational Linguistics.
- Kevin Clark, Minh-Thang Luong, Quoc V Le, and Christopher D Manning. 2019. [Electra: Pre-training text encoders as discriminators rather than generators](#). In *International Conference on Learning Representations*.
- Bert F Green Jr, Alice K Wolf, Carol Chomsky, and Kenneth Laughery. 1961. [Baseball: an automatic question-answerer](#). In *Papers presented at the May 9-11, 1961, western joint IRE-AIEE-ACM computer conference*. ACM, pages 219–224.
- Robin Jia and Percy Liang. 2017. [Adversarial Examples for Evaluating Reading Comprehension Systems](#). In *Proceedings of the 2017 Conference on Empirical Methods in Natural Language Processing*, pages 2021–2031, Copenhagen, Denmark. Association for Computational Linguistics.
- Nelson F. Liu, Roy Schwartz, and Noah A. Smith. 2019. [Inoculation by Fine-Tuning: A Method for Analyzing Challenge Datasets](#). In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 2171–2179, Minneapolis, Minnesota. Association for Computational Linguistics.
- Pranav Rajpurkar, Robin Jia, and Percy Liang. 2018. [Know What You Don’t Know: Unanswerable Questions for SQuAD](#). In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, pages 784–789, Melbourne, Australia. Association for Computational Linguistics.
- Pranav Rajpurkar, Jian Zhang, Konstantin Lopyrev, and Percy Liang. 2016. [SQuAD: 100,000+ Questions for Machine Comprehension of Text](#). In *Proceedings of the 2016 Conference on Empirical Methods in Natural Language Processing*, pages 2383–2392, Austin, Texas. Association for Computational Linguistics.
- Marco Tulio Ribeiro, Tongshuang Wu, Carlos Guestrin, and Sameer Singh. 2020. [Beyond Accuracy: Behavioral Testing of NLP Models with CheckList](#). In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 4902–4912, Online. Association for Computational Linguistics.
- Frances Theisen, Rebecca Leda, Vincent Pozorski, Jennifer M Oh, Nagesh Adluru, Rachel Wong, Ozioma Okonkwo, Douglas C Dean 3rd, Barbara B Bendlin, Sterling C Johnson, Andrew L Alexander, Catherine L Gallagher. 2017. [Evaluation of striatonigral connectivity using probabilistic tractography in Parkinson’s disease](#). *Neuroimage Clin*. 2017 Sep 9;16:557-563. doi: 10.1016/j.nicl.2017.09.009.
- Thomas Wolf, Lysandre Debut, Victor Sanh, Julien Chaumond, Clement Delangue, Anthony Moi, Pierric Cistac, Tim Rault, R’emi Louf, Morgan Funtowicz, and Jamie Brew. 2019. [Huggingface’s transformers: State-of-the-art natural language processing](#). ArXiv, abs/1910.03771.
- Zhilin Yang, Peng Qi, Saizheng Zhang, Yoshua Bengio, William Cohen, Ruslan Salakhutdinov, and Christopher D. Manning. 2018. [HotpotQA: A Dataset for Diverse, Explainable Multi-hop Question Answering](#). In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pages 2369–2380, Brussels, Belgium. Association for Computational Linguistics.

Notebook used to train and work with model(s) for NLP F22 final project

```
from google.colab import drive
drive.mount('/content/drive', force_remount=True)

# change to proper directory
%cd /content/drive/MyDrive/fp-dataset-artifacts
#!git pull origin main
%ls

import os

os.environ['USER'] = 'jackgopack4'
os.environ['PASS'] = 'ghp_hMhEwNAuJqUX92fXuMkXKP6LM1E8HR0kNdX1'
os.environ['REPO'] = 'gregdurrett/fp-dataset-artifacts'
# Clone Prof Durrett's github repo for final project
!git clone https://$USER:$PASS@github.com/$REPO.git # clone the repo
```

Download the release data from Checklist github repository

```
!wget https://github.com/marcotcr/checklist/blob/master/release_data.tar.gz
!tar -xvzf release_data.tar.gz
```

Format the release data from Checklist repo in order to be imported using HuggingFace Datasets

```
import json
def load_squad(fold='validation'):
    answers = []
    data = []
    ids = []
    files = {
        'validation': '/content/drive/MyDrive/fp-dataset-artifacts/release_data/squad/squad.json',
    }
    f = json.load(open(files[fold]))
    i = 0
    for t in f['data']:
        i+=1
        for p in t['paragraphs']:
            i+=1
            context = p['context']
            for qa in p['qas']:
                d = {'id': str(i), 'context': context, 'question': qa['question'], 'answers': {'text': list(["", "", ""]), 'answer_start': list(['0', '0', '0'])}}
                #print(d)
                data.append(d)
                #print(qa['answers'])
                #answers.append(set([(x['text'], x['answer_start']) for x in qa['answers']]))
            with open(os.path.join('/content/drive/MyDrive/fp-dataset-artifacts/release_data/squad/', 'squad_formatted.jsonl'), encoding='utf-8', mode='w') as f2:
                #f.write(json.dumps(data))
                for d in data:
                    f2.write(json.dumps(d))
                    f2.write('\n')
            return data
    load_squad()

# Select "Use fallback runtime version" in Colab Command Palette
# (only available until mid-december), validate 3.7.15 is version shown
!python3 --version

# install the dependencies
!python3 -m pip install --upgrade pip
!python3 -m pip install -r requirements.txt
```

```
#train the model
!python3 run.py --do_train --task qa --model ./trained_model-combined-3 --dataset ./huggingface/huggingface/datasets/experiments/squad_negation_former_latter.jsonl --output_dir ./trained_model-combined-3-negation-former-latter/
```

Evaluate on CheckList squad data (category-based)

```
# evaluate the model
!python3 run.py --do_eval --task qa --dataset /content/drive/MyDrive/fp-dataset-artifacts/release_data/squad/squad_formatted.jsonl --model ./trained_model-combined-3-negation-former-latter/ --output_dir ./eval_output_checklist_combined-3-negation-former-latter/
```

Evaluate on squad data (overall accuracy)

```
!python3 run.py --do_eval --task qa --dataset adversarial_qa --model ./trained_model-combined-3-negation-2/ --output_dir ./eval_output_adversarial_qa-combined-3-negation-2/
```

```
!rm -rf /content/drive/MyDrive/fp-dataset-artifacts/eval_output_og_checklist_full_model/
```

Following code adapted from checklist repo at <https://github.com/marcotcr/checklist>

```
!jupyter nbextension install --py --sys-prefix checklist.viewer
!jupyter nbextension enable --py --sys-prefix checklist.viewer
```

```
#!/pip install checklist
import checklist
from checklist.test_suite import TestSuite
import logging
logging.basicConfig(level=logging.ERROR)
suite_path = '/content/drive/MyDrive/fp-dataset-artifacts/release_data/squad/squad_suite.pkl'
suite = TestSuite.from_file(suite_path)
#print(suite.get_raw_example_List()[0:100])

pred_path = '/content/drive/MyDrive/fp-dataset-artifacts/eval_output_og_checklist_combined/eval_predictions.jsonl'
suite.run_from_file(pred_path, overwrite=True, file_format='pred_only')
suite.summary()
#suite.visual_summary_table()
```

804

## 805 B Exact Match % / F1 Results for all 806 experimental models

```
# Results for squad-trained model:
{'eval_exact_match': 78.01324503311258, 'eval_f1': 86.09164447174685}

# Results for squad combined with AddOneSent trained model:
{'eval_exact_match': 79.91485335856197, 'eval_f1': 87.27578558154276}

# Results for squad with AddSent trained model:
{'eval_exact_match': 80.23651844843897, 'eval_f1': 87.31952701446365}

# combined plus adversarial dataset
{'eval_exact_match': 69.92431409649953, 'eval_f1': 78.54105339971485}

# squad plus adversarial dataset
{'eval_exact_match': 69.2336802270577, 'eval_f1': 78.20658141400816}

# triple dataset
{'eval_exact_match': 80.49195837275307, 'eval_f1': 87.54094922682714}

# triple dataset trained on custom former-latter dataset
{'eval_exact_match': 79.60264900662251, 'eval_f1': 86.38583782036187}

# triple dataset double-trained on custom former-latter dataset
{'eval_exact_match': 79.45127719962157, 'eval_f1': 86.24196078975284}

# squad-trained model with custom former-latter dataset
{'eval_exact_match': 77.22800378429517, 'eval_f1': 84.69200902152039}

# combined-triple dataset with negation experiment set
{'eval_exact_match': 80.33112582781457, 'eval_f1': 87.42694235095531}

# combined-triple dataset with negation-2 experiment set
{'eval_exact_match': 80.0473036896878, 'eval_f1': 86.99738144846238}

# combined triple with negation and former latter experiment set
{'eval_exact_match': 79.46073793755913, 'eval_f1': 86.58063084794101}
```

807